

Intelligent Multimedia Systems

Master AI, 2012, Lecture 4

Lecturers: Theo Gevers

Lab: Intelligent Systems Lab Amsterdam (ISLA)

Email: th.gevers@uva.nl

<http://staff.science.uva.nl/~gevers>




Lectures

- 29-10-2012, Monday, 15:00-17:00, Science Park A1.04 - Introduction
- 05-11-2011, Monday, 15:00-17:00, Science Park A1.04 - Image and Video Formation
- 12-11-2011, Monday, 15:00-17:00, Science Park A1.04 - Color Invariance and Image Processing
- 19-11-2011, Monday, 15:00-17:00, Science Park A1.04 - Feature Extraction and Tracking
- 26-11-2011, Monday, 15:00-17:00, Science Park A1.04 - Learning and Object Recognition
- 03-12-2011, Monday, 15:00-17:00, Science Park A1.04 - Visual Attention and Affective Computing
- 10-12-2011, Monday, 15:00-17:00, Science Park A1.04 - Human Behavior Analysis
- 18-12-2011, Tuesday, 15:00-18:00, Science Park, C1.10 - Examination

Today's class

Feature Extraction

Tracking




Local Features: Main Components

1) Detection: Identify the interest points

2) Description: Extract vector feature descriptor surrounding each interest point.


$$\mathbf{x}_1 = [x_1^{(1)}, \dots, x_d^{(1)}]$$



$$\mathbf{x}_2 = [x_1^{(2)}, \dots, x_d^{(2)}]$$

3) Matching: Determine correspondence between descriptors in two images

Feature-based Alignment



- Extract features
- Compute *putative matches*
- Loop:
 - *Hypothesize* transformation T (small group of putative matches that are related by T)
 - *Verify* transformation (search for other matches consistent with T)

Feature-based Alignment



- Extract features
- Compute *putative matches*
- Loop:
 - *Hypothesize* transformation T (small group of putative matches that are related by T)
 - *Verify* transformation (search for other matches consistent with T)

Automatic Mosaicing



<http://www.cs.ubc.ca/~mbrown/autostitch/autostitch.html>

Recognition of Specific Objects, Scenes



Schmid and Mohr 1997



Sivic and Zisserman, 2003



Rothganger et al. 2003



Lowe 2002


Preview

- Interest point detection
 - Edges: derivatives of Gaussians (first and second order)
 - Blobs: Laplacian of Gaussian, automatic scale selection
 - Harris corner detector
 - Template matching
- Invariant descriptors
 - Rotation according to dominant gradient direction
 - Histograms for robustness to small shifts and translations (SIFT descriptor)

Feature Extraction

by



Image Filtering




Recap: Edge Detection

1	0	-1
2	0	-2
1	0	-1

intensity image




Recap: Edge Detection




Recap: Edge Detection

f




Edge

$\frac{d^2}{dx^2} g$



Second derivative
of Gaussian
(Laplacian)




$f * \frac{d^2}{dx^2} g$



Edge = zero crossing
of second derivative

Summary

- Interest point detection
 - Edges: derivatives of Gaussians (first and second order)
 - Blobs: Laplacian of Gaussian, automatic scale selection



Classifying of Color Edges



Edge Detection



An intensity edge is defined as a point where...

the gradient is large:


$$|\nabla f(x, y)| = \sqrt{f_x^2 + f_y^2} \gg 0$$

Localization - Laplacian - zero-crossing:

$$\nabla^2 f(x, y) = f_{xx} + f_{yy}$$



gradient computation



zero-crossings at large gradients

Edge Detection

Summing the gradient magnitudes separately:

$$|\nabla C(x, y)| = \sqrt{(R_x^2 + R_y^2)} + \sqrt{(G_x^2 + G_y^2)} + \sqrt{(B_x^2 + B_y^2)}$$

or using the Euclidean metric:

$$|\nabla C(x, y)| = \sqrt{R_x^2 + R_y^2 + G_x^2 + G_y^2 + B_x^2 + B_y^2}$$

or using eigen-values: [diZenzo86], [Sapiro96]

Edge Localization

Edge localization by:

1. Non-maxima suppression based on the direction of the minimal and maximal change given by the eigen-vectors:

$$\theta_- \text{ and } \theta_+$$

2. Zero-crossing (change in sign) by scanning:



$$\nabla^2 C_{RGB} = (R_{xx} + R_{yy} + G_{xx} + G_{yy} + B_{xx} + B_{yy})$$

$$\nabla^2 C_{c_1 c_2 c_3} = (c_{1xx} + c_{1yy} + c_{2xx} + c_{2yy} + c_{3xx} + c_{3yy})$$


$$\nabla^2 C_{l_1 l_2 l_3} = (l_{1xx} + l_{1yy} + l_{2xx} + l_{2yy} + l_{3xx} + l_{3yy})$$

Edge Classification


summing
gradient
magnitudes
separately:



RGB





$C_1 C_2 C_3$




$I_1 I_2 I_3$


Euclidean:



RGB




$C_1 C_2 C_3$




$I_1 I_2 I_3$


eigen-values:



RGB





$C_1 C_2 C_3$




$I_1 I_2 I_3$


Edge Classification



RGB



c_1, c_2, c_3



l_1, l_2, l_3



if $(|\nabla C_{c_1 c_2 c_3}| \geq t_{c_1 c_2 c_3} \& |\nabla C_{l_1 l_2 l_3}| < t_{l_1 l_2 l_3})$ then classify as highlight edge

else


if $(|\nabla C_{l_1 l_2 l_3}| \geq t_{l_1 l_2 l_3})$ then classify as color edge

else classify as shadow/geometry edge


Edge Classification




colour edge maxima by type



shadows and geometry



highlights




colour edges


Edge Classification

material

shadow or geometry




Edge Classification




material
highlight
shadow or geometry

Edge Classification



material
highlight
shadow or geometry



Demo: Edge Classification


video



classification



shadow-shading



Demo: Edge Classification

video



classification



material



shadow-shading

Deformable Contours

Contour which minimises elastic energy:



$$b|v_k(s)|^2 + f(v(s))$$



Advanced segmentation suited for query formulation

Deformable Contours

Deformable curve : $v(t) = [x(t), y(t)], t \in [0,1]$

The energy : $E = \alpha E_{\text{int}} + \beta E_{\text{ext}}$


Internal energy : $E_{\text{int}} = \left(\oint_t (\|v(t)'\|^2 + \|v(t)''\|^2) dt \right) \left(\oint_t \|v(t)'\|^2 \right)$

External energy (intensity gradient) : $E_{\text{ext}} = \oint_t -\|\nabla I(x, y)\| dt$


External energy (color invariant gradient) : $E_{\text{ext}} = \oint_t -\|\nabla C(x, y)\| dt$




Deformable Contours



Deformable Contours



Demo: Tracking by Deformable Contours



Demo: Tracking by Deformable Contours





Shadow Removal

by

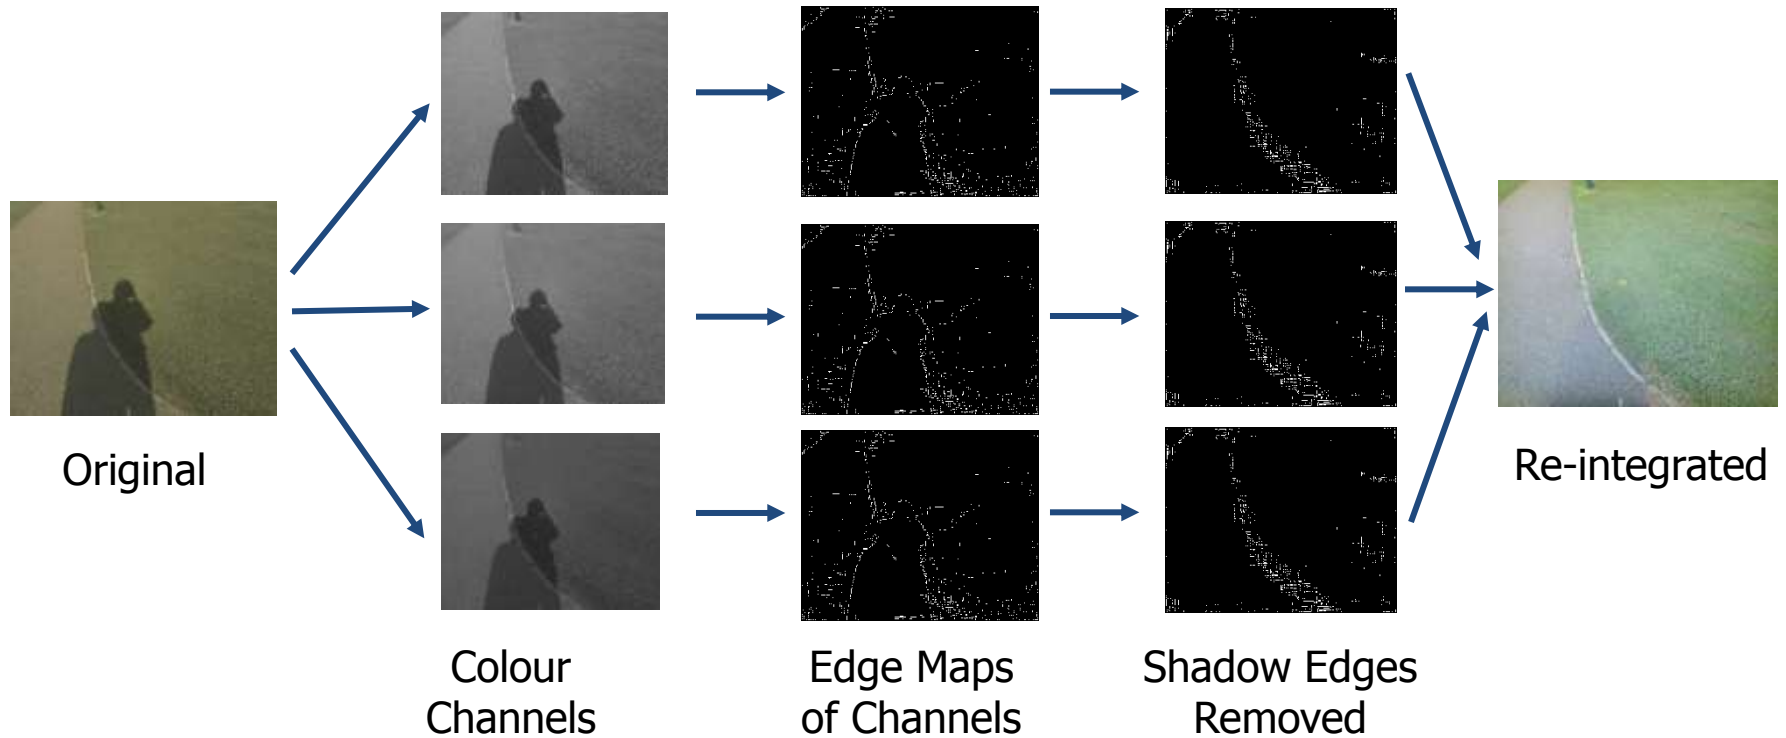
Graham Finlayson

Shadow Removal



We would like to go from a colour image with shadows, to the same colour image, but without the shadows.

Shadow Removal



An Example

Original
Image



Invariant
Image



Detected
Shadow Edges



Shadow
Removed



A Second Example


Original
Image



Invariant
Image



Detected
Shadow Edge



Shadow
Removed



More Examples


Original
Image



Invariant
Image




Detected
Shadow Edges



Shadow
Removed



Intrinsic Images




Intrinsic Images




Feature Extraction


by

Template Matching



Intermezzo: Template Matching

- Goal: find  in image
- Main challenge: What is a good similarity or distance measure between two patches?
 - Zero-mean correlation
 - Sum Square Difference
 - Normalized Cross Correlation




Intermezzo: Template Matching

- Goal: find  in image
- Method 1: filter the image with zero-mean eye

$$h[m,n] = \sum_{k,l} (f[k,l] - \bar{f})(g[m+k,n+l])$$


\bar{f} ← mean of f



Input




Filtered Image (scaled)




Thresholded Image

Intermezzo: Template Matching

- Goal: find  in image
- Method 2: SSD


$$h[m,n] = \sum_{k,l} (g[k,l] - f[m+k,n+l])^2$$



Input



1- sqrt(SSD)



Thresholded Image


Intermezzo: Template Matching


- Goal: find  in image
- Method 3: Normalized cross-correlation

$$h[m,n] = \frac{\sum_{k,l} (g[k,l] - \overset{\text{mean template}}{\downarrow} \bar{g})(f[m-k,n-l] - \overset{\text{mean image patch}}{\downarrow} \bar{f}_{m,n})}{\left(\sum_{k,l} (g[k,l] - \bar{g})^2 \sum_{k,l} (f[m-k,n-l] - \bar{f}_{m,n})^2 \right)^{0.5}}$$

Matlab: `normxcorr2(template, im)`

Intermezzo: Template Matching


- Goal: find  in image
- Method 3: Normalized cross-correlation



Input



Normalized X-Correlation



Thresholded Image

What is the best method to use?

A: Depends

- SSD: faster, sensitive to overall intensity
- Normalized cross-correlation: slower, invariant to local average intensity and contrast

Summary


- Applications of filters
 - Template matching (SSD or Normxcorr2)
 - SSD can be done with linear filters, is sensitive to overall intensity
 - Gaussian pyramid
 - Coarse-to-fine search, multi-scale detection
 - Downsampling
 - Need to sufficiently low-pass before downsampling

Feature Extraction

by


Image Filtering:

Corners




An introductory example:

Harris corner detector




The Basic Idea


- We should easily recognize the point by looking through a small window
- Shifting a window in *any direction* should give a *large change* in intensity




Harris Detector: Basic Idea



“flat” region:
no change in
all directions



“edge”:
no change along
the edge direction



“corner”:
significant change
in all directions

Harris Detector: Mathematics

Change of intensity for the shift $[u, v]$:


$$E(u, v) = \sum_{x, y} w(x, y) [I(x + u, y + v) - I(x, y)]^2$$

Window
function

Shifted
intensity


Intensity

Window function $w(x, y) =$



1 in window, 0 outside

or



Gaussian

Harris Detector: Mathematics

For small shifts $[u, v]$ we have a *bilinear* approximation:

$$E(u, v) \cong [u, v] M \begin{bmatrix} u \\ v \end{bmatrix}$$

where M is a 2x2 matrix computed from image derivatives:


$$M = \sum_{x, y} w(x, y) \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

Harris Detector: Mathematics

Intensity change in shifting window: eigenvalue analysis

$$E(u, v) \cong [u, v] M \begin{bmatrix} u \\ v \end{bmatrix} \quad \lambda_1 \text{ and } \lambda_2 - \text{eigenvalues of } M$$


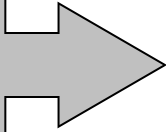
Ellipse $E(u, v) = \text{const}$



Harris Detector: Mathematics

Classification of image points using eigenvalues of M :

λ_1 and λ_2 are small;
 E is almost constant
in all directions



Harris Detector: Mathematics

Measure of corner response:

$$R = \det M - k (\text{trace } M)^2$$

$$\det M = \lambda_1 \lambda_2$$

$$\text{trace } M = \lambda_1 + \lambda_2$$

(k – empirical constant, $k = 0.04$ - 0.06)

Summary of the Harris detector

1. Compute x and y derivatives of image

$$I_x = G_\sigma^x * I \quad I_y = G_\sigma^y * I$$

2. Compute products of derivatives at every pixel

$$I_{x2} = I_x \cdot I_x \quad I_{y2} = I_y \cdot I_y \quad I_{xy} = I_x \cdot I_y$$

3. Compute the sums of the products of derivatives at each pixel

$$S_{x2} = G_{\sigma^2} * I_{x2} \quad S_{y2} = G_{\sigma^2} * I_{y2} \quad S_{xy} = G_{\sigma^2} * I_{xy}$$

4. Define at each pixel (x, y) the matrix

$$H(x, y) = \begin{bmatrix} S_{x2}(x, y) & S_{xy}(x, y) \\ S_{xy}(x, y) & S_{y2}(x, y) \end{bmatrix}$$

5. Compute the response of the detector at each pixel

$$R = \text{Det}(H) - k(\text{Trace}(H))^2$$


6. Threshold on value of R . Compute nonmax suppression.

Harris Detector [Harris88]

- Second moment matrix (autocorrelation matrix)

$$\mu(\sigma_I, \sigma_D) = g(\sigma_I) * \begin{bmatrix} I_x^2(\sigma_D) & I_x I_y(\sigma_D) \\ I_x I_y(\sigma_D) & I_y^2(\sigma_D) \end{bmatrix}$$

1. Image derivatives



2. Square of derivatives



3. Gaussian filter $g(\sigma_I)$




- 4. Cornerness function – both eigenvalues are strong

$$har = \det[\mu(\sigma_I, \sigma_D)] - \alpha [\text{trace}(\mu(\sigma_I, \sigma_D))^2] =$$


$$g(I_x^2)g(I_y^2) - [g(I_x I_y)]^2 - \alpha [g(I_x^2) + g(I_y^2)]^2$$

- 5. Non-maxima suppression



Harris Detector: Mathematics

- R depends only on eigenvalues of M
- R is large for a corner
- R is negative with large magnitude for an edge
- $|R|$ is small for a flat region



Harris Detector


- The Algorithm:
 - Find points with large corner response function R ($R > \text{threshold}$)
 - Take the points of local maxima of R

Harris Detector: Workflow




Harris Detector: Workflow

Compute corner response R



Harris Detector: Workflow

Find points with large corner response: $R > \text{threshold}$



Harris Detector: Workflow

Take only the points of local maxima of R




Harris Detector: Workflow




Feature selection

- Distribute points evenly over the image




Adaptive Non-maximal Suppression

- Desired: Fixed # of features per image
 - Want evenly distributed spatially...
 - Sort points by non-maximal suppression radius
[Brown, Szeliski, Winder, CVPR'05]



(a) Strongest 250



(b) Strongest 500




(c) ANMS 250, $r = 24$



(d) ANMS 500, $r = 16$

Harris Detector: Some Properties

- Rotation invariance



Ellipse rotates but its shape (i.e. eigenvalues) remains the same

Corner response R is invariant to image rotation

Harris Detector: Some Properties

- Partial invariance to *affine intensity* change

Only derivatives are used => invariance to intensity shift $I = I + b$



Intensity scale: $I = a I$

Harris Detector: Some Properties

- Quality of Harris detector for different scale changes

Repeatability rate:

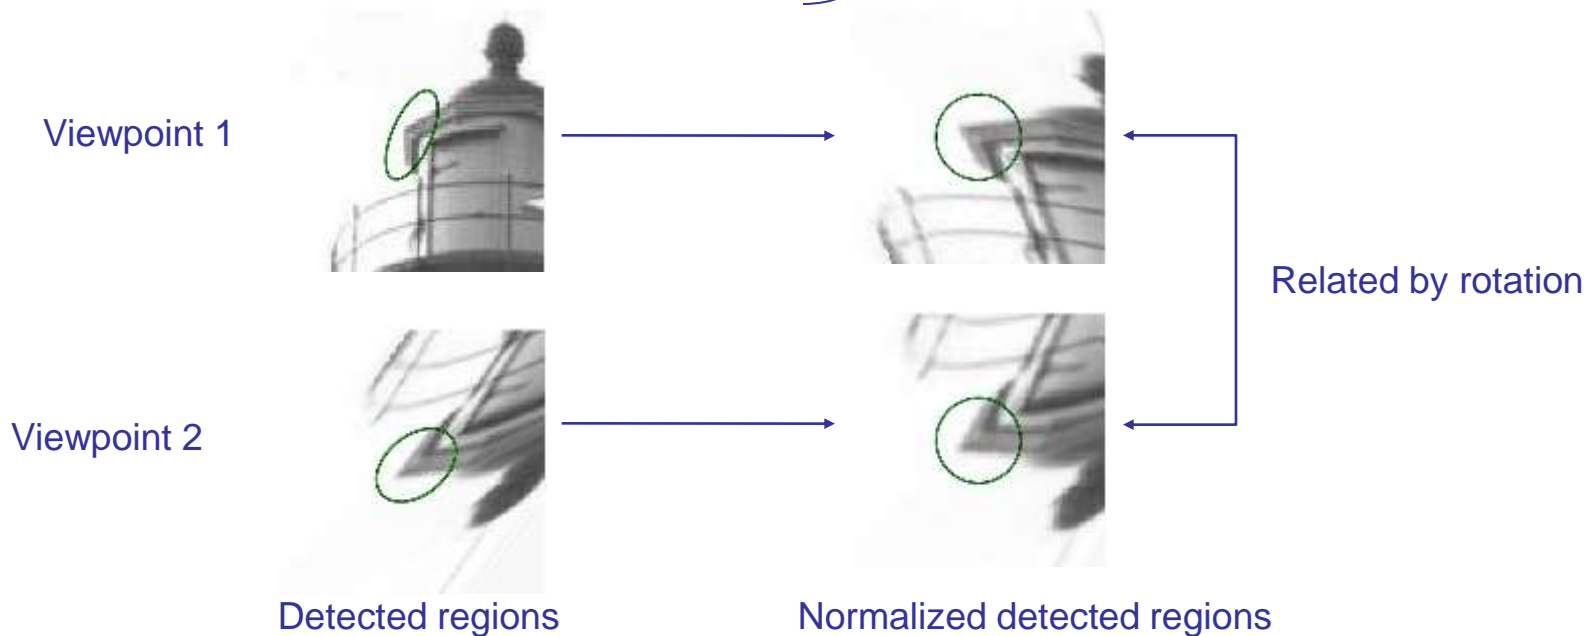
$$\frac{\# \text{ correspondences}}{\# \text{ possible correspondences}}$$



Matching with Features

- Rotation
- Scaling
- Viewpoint

Affine invariance
!




Matching with Features

- Existing method by Mikolajczyk
 - Iterative affine invariant point detector
 - Multi-scale Harris corner detector
 - Laplacian characteristic scale selection
 - Second moment matrix shape determination



Initial region based on initial scale and location



Iteratively adjust scale, position and shape of region




final region



Original Image



Harris Laplacian
impl. by Mikolajczyk (e.g. CVPR06)



Shape adapted Harris Laplacian
impl. by Mikolajczyk (ICCV07)



Color salient points
Quasi invariant HSI



Color salient points
Color boosted OCS

Most of the time, both color approaches agree on the most salient parts of an image.

Comparison of Keypoint Detectors


Table 7.1 Overview of feature detectors.

Feature Detector	Corner	Blob	Region	Rotation invariant	Scale invariant	Affine invariant	Repeatability	Localization accuracy	Robustness	Efficiency
Harris	✓			✓			+++	+++	+++	++
Hessian		✓		✓			++	++	++	+
SUSAN	✓			✓			++	++	++	+++
Harris-Laplace	✓	(✓)		✓	✓		+++	+++	++	+
Hessian-Laplace	(✓)	✓		✓	✓		+++	+++	+++	+
DoG	(✓)	✓		✓	✓		++	++	++	++
SURF	(✓)	✓		✓	✓		++	++	++	+++
Harris-Affine	✓	(✓)		✓	✓	✓	+++	+++	++	++
Hessian-Affine	(✓)	✓		✓	✓	✓	+++	+++	+++	++
Salient Regions	(✓)	✓		✓	✓	(✓)	+	+	++	+
Edge-based	✓			✓	✓	✓	+++	+++	+	+
MSER			✓	✓	✓	✓	+++	+++	++	+++
Intensity-based			✓	✓	✓	✓	++	++	++	++
Superpixels			✓	✓	(✓)	(✓)	+	+	+	+

Summary


- Interest point detection
 - Harris corner detector
 - Laplacian of Gaussian, automatic scale selection
- Invariant descriptors
 - Rotation according to dominant gradient direction
 - Histograms for robustness to small shifts and translations (SIFT descriptor)

Local Features






Local features: main components

- 1) Detection: Identify the interest points
- 2) Description: Extract vector feature descriptor surrounding each interest point.
- 3) Matching: Determine correspondence between descriptors in two views



Geometric transformations



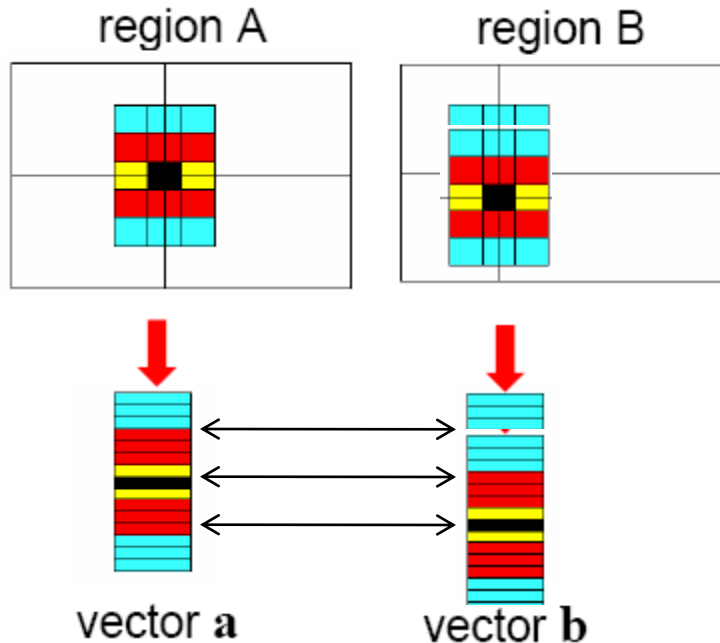
e.g. scale,
translation,
rotation

Photometric transformations



Figure from T. Tuytelaars ECCV 2006 tutorial

Raw patches as local descriptors





The simplest way to describe the neighborhood around an interest point is to write down the list of intensities to form a feature vector.

But this is very sensitive to even small shifts, rotations.


SIFT descriptor [Lowe 2004]

- Use histograms to bin pixels within sub-patches according to their orientation.



*Why subpatches?
Why does SIFT
have some
illumination
invariance?*

Making descriptor rotation invariant



- Rotate patch according to its dominant gradient orientation
- This puts the patches into a canonical orientation.

SIFT descriptor [Lowe 2004]

- Extraordinarily robust matching technique
 - Can handle changes in viewpoint
 - Up to about 60 degree out of plane rotation
 - Can handle significant changes in illumination
 - Sometimes even day vs. night (below)
 - Fast and efficient—can run in real time
 - Lots of code available
 - http://people.csail.mit.edu/albert/ladypack/wiki/index.php/Known_implementations_of_SIFT



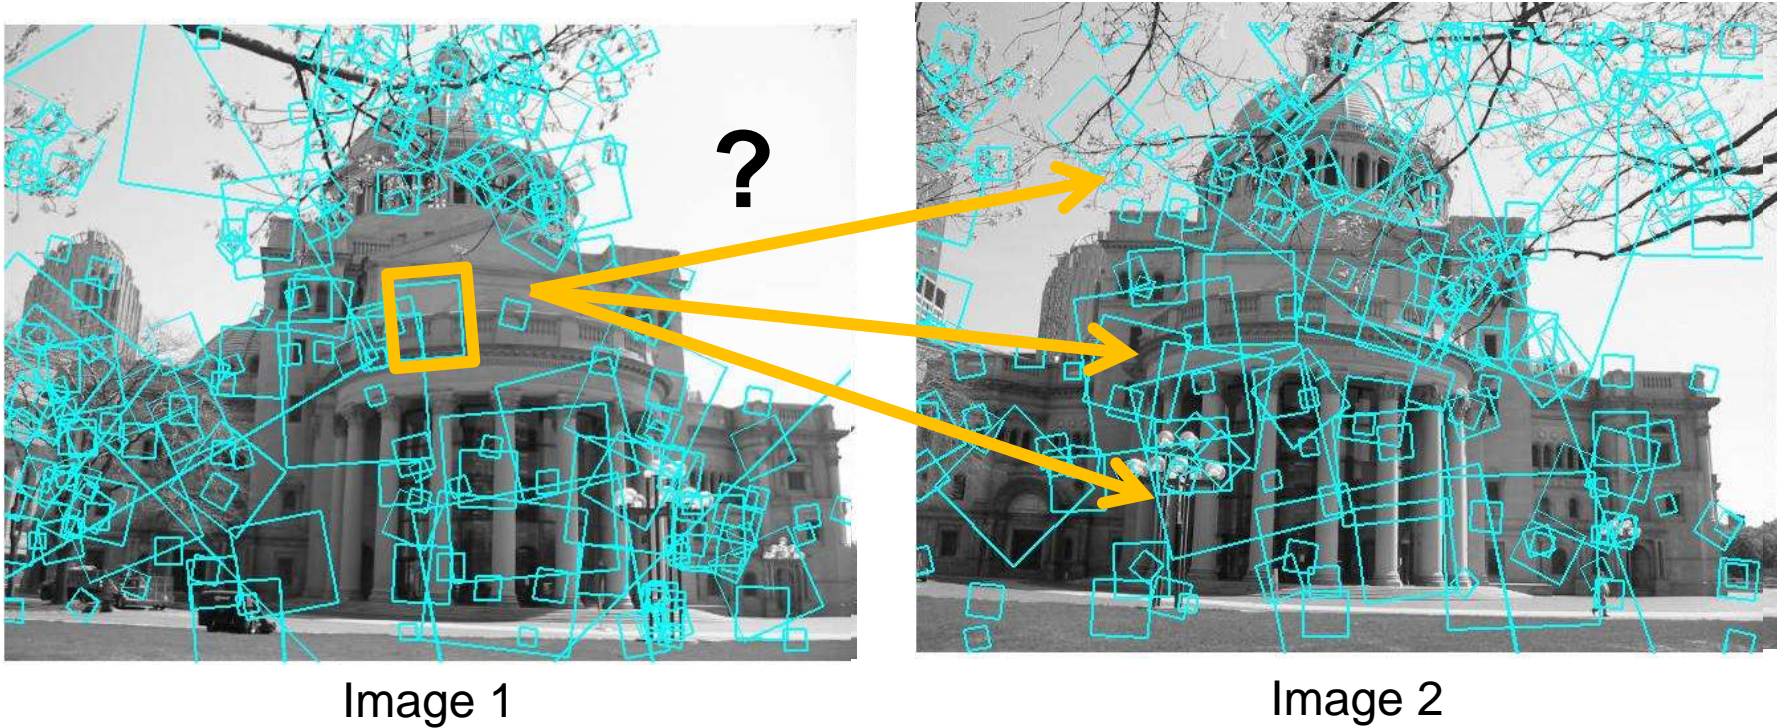
SIFT properties

- Invariant to
 - Scale
 - Rotation
- Partially invariant to
 - Illumination changes
 - Camera viewpoint
 - Occlusion, clutter

Matching local features



Matching local features



To generate **candidate matches**, find patches that have the most similar appearance (e.g., lowest SSD)

Simplest approach: compare them all, take the closest (or closest k , or within a thresholded distance)

Ambiguous matches



Image 1



Image 2

At what SSD value do we have a good match?


To add robustness to matching, can consider **ratio** :
distance to best match / distance to second best match

If low, first match looks good.

If high, could be ambiguous match.

Matching SIFT Descriptors

- Nearest neighbor (Euclidean distance)
- !
- !



Recognition of specific objects, scenes



Schmid and Mohr 1997



Sivic and Zisserman, 2003



Rothganger et al. 2003




Lowe 2002

Today's class

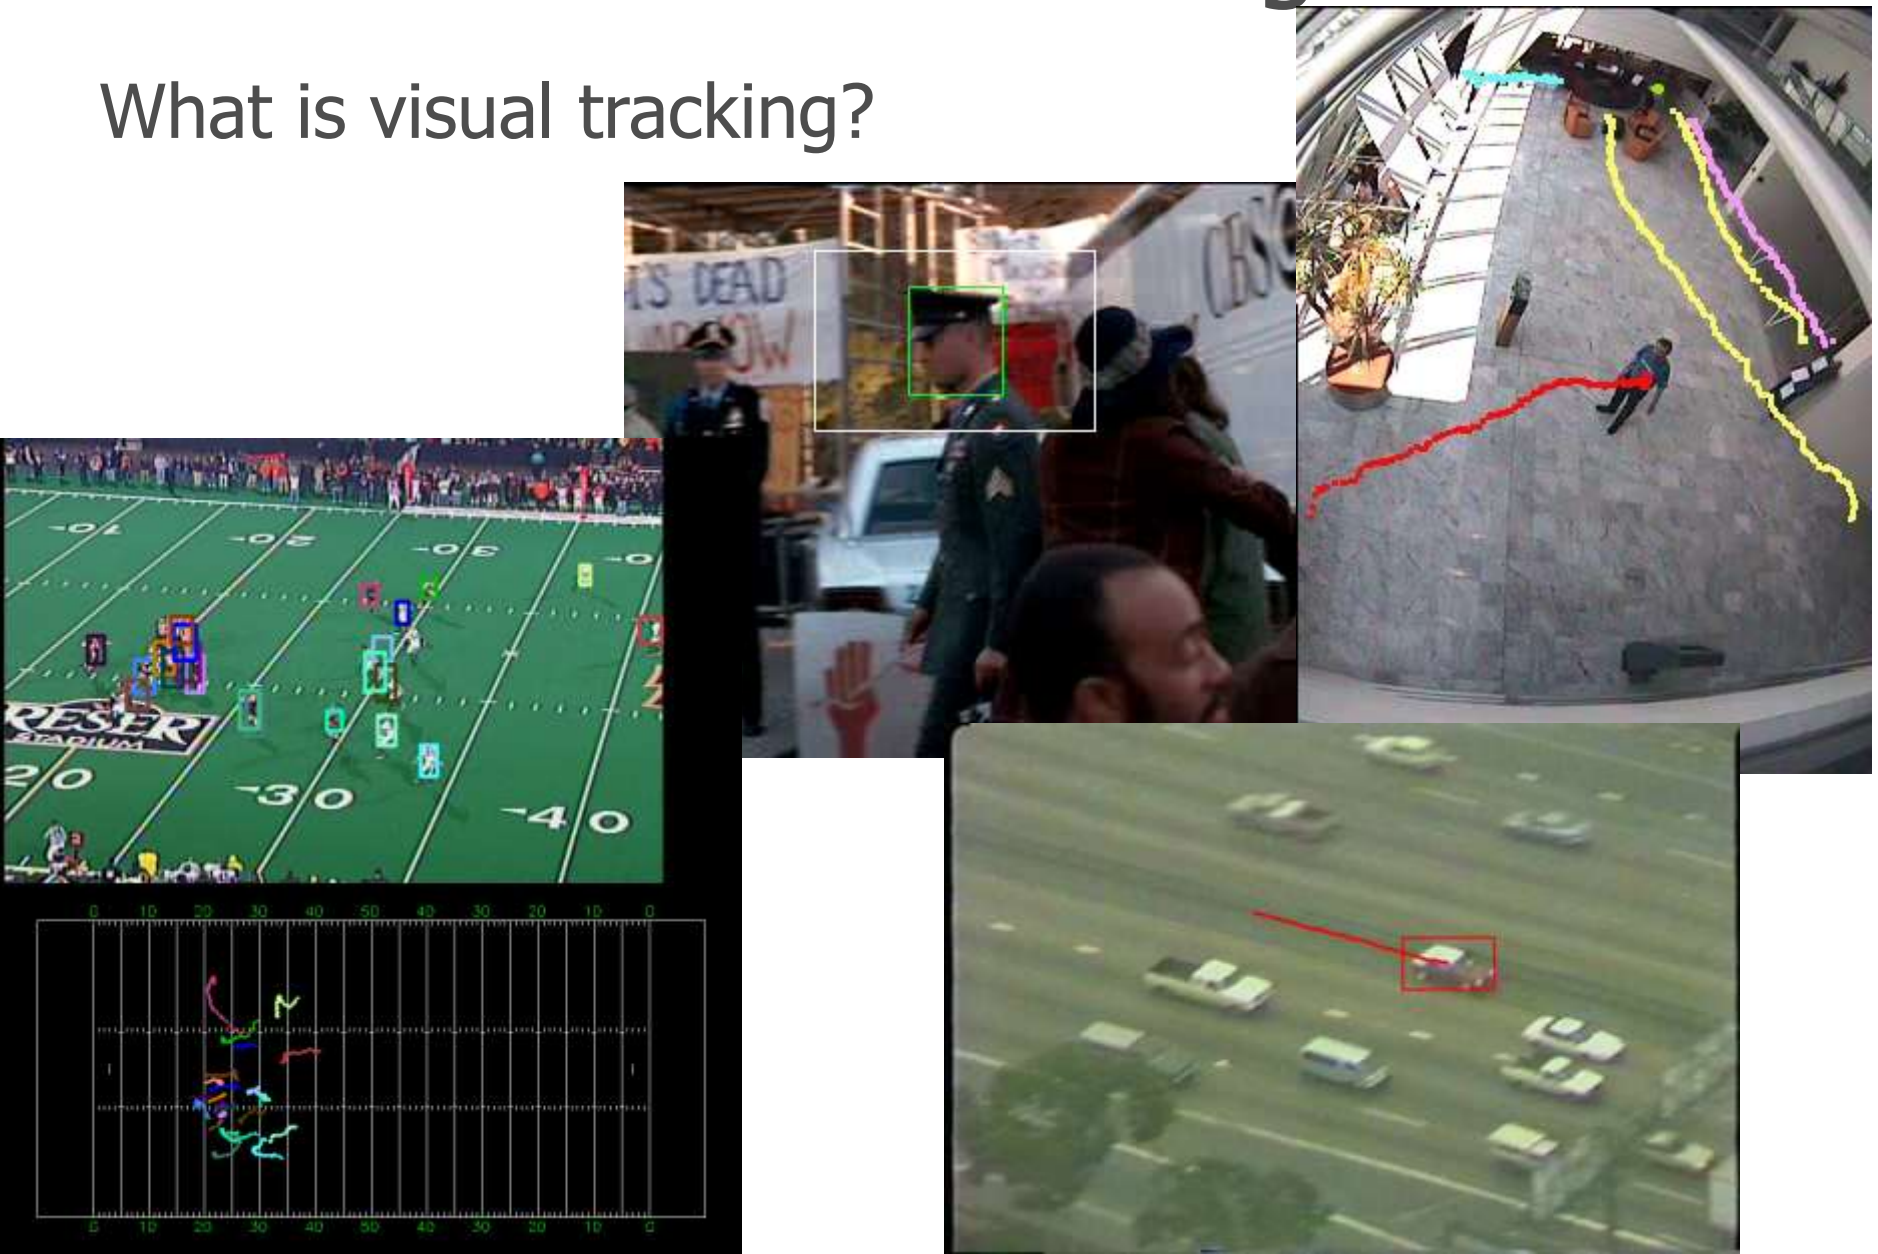
Feature Extraction

Tracking




Visual tracking

What is visual tracking?



Visual tracking


What are the applications of tracking ?



Visual tracking

Tracking can be very easy!

- Tracking can be very easy if both the target and the background are uniform in color.



Visual tracking

What makes tracking so hard?

- **Background clutter:** the presence of other objects or non-informative patterns in the image complicates the detection of the right object.
- A **dynamic background:** moving camera, viewpoint change.
- **Illumination change:** change in direction or intensity of light source, shadow...
- **Non-rigid, unhomogeneous, fast objects:** articulated objects, appearance change, object speed / frame rate
- **Multiple objects tracking**
- **Occlusion:** the target disappears partially or completely behind another object for a while.
- **Setting and application**

Visual tracking

Standard tracking algorithms

- **Motion segmentation based tracking**
- **Template tracking**
- **Mean-shift tracking**
- **Kalman filter**
- **Particle filtering**

Visual tracking

Motion segmentation

Motion segmentation methods aims at separating foreground from background:

- Detections without background model are fast but noisy.

Image difference, Optical flow



Visual tracking

Motion segmentation

- Pixel-based background model

Image model, statistical model (Gaussian mixture model), predictive model

Frame



RGB



YUV



Visual tracking

Motion segmentation



- Local background model
Region detection, texture on block, improve pixel-based model
- Global background model
Model mixture, eigen background

Visual tracking

Tracking methods

Region-based tracking

Connected regions are detected and matched across frames

- Blobs composing human body.

- Various levels tracking: regions, person, group.

Method that handles various object tracking. However occlusions, objects interactions are not handled.

Contour-based tracking

Active contour uses the objects outline.

The representation is simple, but the precision is limited to the contour and initialization is challenging.

Object Tracking: A Survey, ACM Computing Surveys, Yilmaz A., Javed, Shah, 2006.

Visual tracking

Tracking methods

Feature-based tracking

Extract feature to describe and match regions

–Global features: centroid, perimeter, surface area, color moments

–Local features: interest points, corner, curve/lines segments

Multiple object can be tracked and partial occlusion can be handled.

These features can be easily combined.

Model-based tracking

Match a projected object model to image data.

A priori knowledge about the object is required.

Human models include stick figure, volumetric models, skeleton, etc.

Need to build the model. High computational costs.

Visual tracking


Standard tracking algorithms

- **Motion segmentation based tracking**
- **Template tracking**
- **Mean-shift tracking**
- **Kalman filter**
- **Particle filtering**



Visual tracking

Template based tracking

- Tracking consists in searching for the target object in a frame by comparing with a **template** image.
- We assume that the template is fixed and given in advance.



Template image
 $T(\mathbf{x})$




$I(\mathbf{x}, \mathbf{t})$

Visual tracking

Template to target transformation

- The template is mapped into a candidate target region the image using a transformation of coordinates: $\phi(\mathbf{x}): \Omega \rightarrow \square$. This transformation depends on a parameter vector \mathbf{y} . Different candidate regions correspond to different values of \mathbf{y} . So we write $\phi(\mathbf{x}; \mathbf{y})$.



Visual tracking


Motion models

- The type of transformation φ specifies the type of object motion that the tracker is able to deal with.
 - Translation: $\varphi(\mathbf{x}; \mathbf{y}) = \mathbf{x} + \mathbf{y}$
 - Rotation:
 $\varphi_1 = x_1 \cos y - x_2 \sin y$
 $\varphi_2 = x_1 \sin y + x_2 \cos y$
 - Scaling:
 $\varphi_1 = yx_1$
 $\varphi_2 = yx_2$
 - Affine:
 $\varphi_1 = y_1 + y_2x_1 + y_3x_2$
 $\varphi_2 = y_4 + y_5x_1 + y_6x_2$

Visual tracking

Search


- Align the template with every possible candidate region in the image, and find the most similar candidate according to a **similarity measure**.
- We search the target only in an area around the previous position exploiting general knowledge that the object won't have moved far.



Visual tracking

Similarity measure

- We need a measure of how similar (or far apart) the template and the candidate are.



- The similarity measure can be based on:
 - pixelwise intensity (color) difference: **SSD** and **correlation** trackers,
 - histogram difference: **mean-shift** tracker.

Visual tracking

SSD and correlation

- SSD is short for sum-of-squared-difference:

$$D(\mathbf{y}) = \sum_{\mathbf{x} \in \Omega} [I(\mathbf{x} + \mathbf{y}) - T(\mathbf{x})]^2 \rightarrow \min_{\mathbf{y}}$$

- A simpler similarity measure is the (unnormalized) cross-correlation:

$$C(\mathbf{y}) = \sum_{\mathbf{x} \in \Omega} I(\mathbf{x} + \mathbf{y})T(\mathbf{x}) \rightarrow \max_{\mathbf{y}}$$

Visual tracking

Exhaustive search

- Calculate SSD for every y in a search window and choose the position with the least SSD.
- Strengths: robustness and simplicity in implementation.
- Weaknesses:
 - Computations could be time-consuming in case of a large search window.
 - Only suitable for translation.


Visual tracking

Gradient descent

- Iteratively update the solution estimate using Newton's method:

$$\mathbf{y}_{n+1} = \mathbf{y}_n - \alpha \nabla D$$

- Efficient computation. Able to cope with rotation.
- Finds a local minimum only.



Visual tracking

Standard tracking algorithms

- **Motion segmentation based tracking**
- **Template tracking**
- **Mean-shift tracking**
- **Kalman filter**
- **Particle filtering**

Visual tracking

Mean-shift tracking

- Target detection is performed by matching **weighted histograms**.
- Very fast in comparison with SSD or correlation trackers.

***Real time tracking of Non-Rigid Objects using Mean Shift,
CVPR, Comaniciu et al. 2000.***

Visual tracking


The mean-shift algorithm

- The mean-shift algorithm finds a local maximum of a density function of the form:


$$f(\mathbf{y}) = \sum_i w_i K\left(\frac{|\mathbf{y} - \mathbf{x}_i|^2}{\sigma}\right)$$

- where K is the local kernel.

Gaussian kernel:



$$K(|\mathbf{x}|^2) = (2\pi)^{-d/2} \exp(-|\mathbf{x}|^2 / 2)$$




Visual tracking

The mean-shift algorithm

- A local maximum will be found by successively shifting \mathbf{y} to a weighted mean of \mathbf{x}_i computed with the *derivative kernel* K' :


$$\mathbf{y}_{n+1} = \frac{\sum_{i=1}^N \mathbf{x}_i w_i K'_i \left(\frac{|\mathbf{y}_n - \mathbf{x}|^2}{\sigma} \right)}{\sum_{i=1}^N w_i K'_i \left(\frac{|\mathbf{y}_n - \mathbf{x}|^2}{\sigma} \right)}$$



Visual tracking

Similarity measure

- $P(i)$: the template histogram,
- $Q(i;\mathbf{y})$: the histogram of the test region,



- *The Bhattacharyya coefficient* can measure the similarity between two distributions:

$$r(\mathbf{y}) = r(P, Q(\mathbf{y})) = \sum_{i=0}^{255} \sqrt{P(i)Q(i;\mathbf{y})} \rightarrow \max_{\mathbf{y}}$$

Visual tracking

Mean-shift tracking

- Tracking under occlusions
- Updating process

Visual tracking

Standard tracking algorithms

- **Motion segmentation based tracking**
- **Template tracking**
- **Mean-shift tracking**
- **Kalman filter**
- **Particle filtering**

Visual tracking

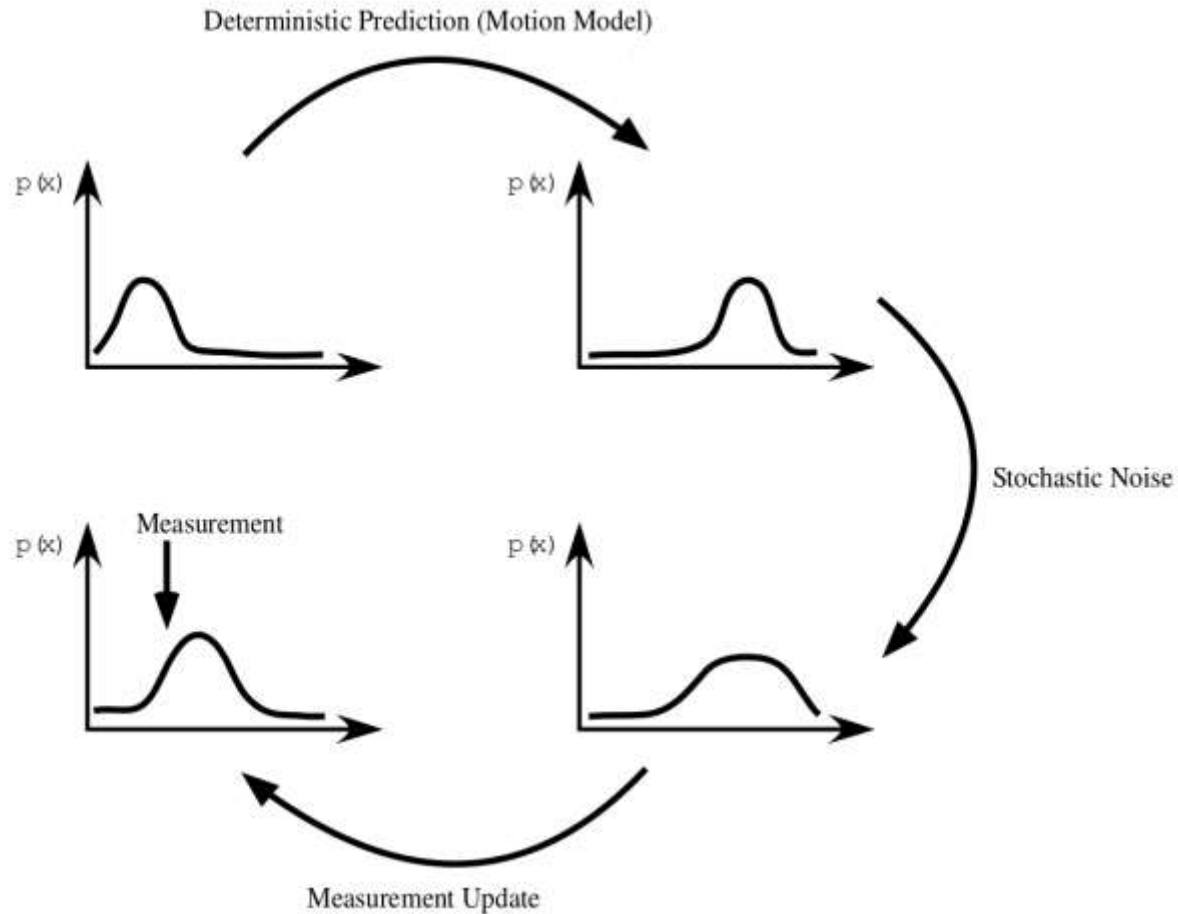
Kalman Filtering

- Provides an optimal estimation of the state of system based on measurements
- Assumptions: linear relationship between variables, distributions are Gaussian (mean, variance).
- 3 steps:
 - Deterministic prediction
 - Stochastic diffusion
 - Estimate refinement via a new measurement

A Review of Visual Tracking, tech. report, Cannons, 2008

Visual tracking

Kalman Filtering



Visual tracking

Kalman Filtering

- x_t state of the system at time t and z_t measurements

$$\hat{x}_t^- = A\hat{x}_{t-1} + B\hat{u}_{t-1}$$

- A : system equation (link state at time t-1 to state at time t)
- B : control input: u is omitted
- Simple example: constant velocity

$$x_t = \{p_x, p_y, v_x, v_y\}^T$$

$$A = \begin{bmatrix} 1 & 0 & 1 & 0 \\ 0 & 1 & 0 & 1 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

Visual tracking

Kalman Filtering

- \hat{P}_t^- covariance estimate $\hat{P}_t^- = AP_{t-1}A^T + Q$
- Q : covariance of the noise associated with the prediction process
- Posterior estimate $\hat{x}_t = \hat{x}_t^- + K_t(z_t - H\hat{x}_t^-)$
- With $K_t = P_t^- H^T (HP_t^- H^T + R)^{-1}$
- R : noise associated with the measurement process
- H : maps state vectors into measurements vector (Identity in simple cases)
- Posterior estimate $P_t = (I - K_t H)P_t^-$

Visual tracking

Kalman Filtering

- K_t Kalman Gain: determines how much the prediction is trusted against the measurements
- Kalman advantages: combines multiple sources of information in an optimal manner, recursiveness (efficient computation).
- Kalman disadvantages: can not model all systems (linear), requires Gaussian models (problem with cluttered: find one optimal solution)

Visual tracking

Standard tracking algorithms

- **Motion segmentation based tracking**
- **Template tracking**
- **Mean-shift tracking**
- **Kalman filter**
- **Particle filtering**


Visual tracking

Particle Filtering

- Origin comes from monte carlo: use random sampling to find an approximated solution
- Condensation algorithm: no need to be linear, handle multimodal distribution
- Unlike Kalman samples have unspecified shape
- 3 steps:
 - Deterministic prediction
 - Stochastic diffusion
 - Correction via measurements updates

Visual tracking

Particle Filtering



Visual tracking

Particle Filtering

- The goal is to compute $p(x_t | z_t)$ probability distribution that describes the system state, given observed measurements

$$p(x_t | z_t) = k \ p(z_t | x_t) p(x_t)$$

- k constant. A set of sample $\{s_t^{(1)} \dots s_t^{(N)}\}$ is drawn from the prior distribution $p(x_t)$
- Then samples are weighted

$$\pi_t^n = \frac{p(z_t | x_t = s_t^{(n)})}{\sum_{j=1}^N p(z_t | x_t = s_t^{(j)})}$$

- Larger N = better precision = more computation

Visual tracking

Particle Filtering

- Step1: sample from approximated posterior $\{(s_{t-1}^{(n)}, \pi_{t-1}^{(n)}), n = 1, \dots, N\}$

Deterministic drift

- Step 2: Add noise to the modified samples to obtain an estimate of

$$p(x_t | z_{t-1})$$

- Step 3: Determines the samples weights (Bhattacharrya: distance between target and candidates)

Visual tracking

Particle Filtering


- Advantages over Kalman:
 - Can track through clutter
 - No need to calculate covariance estimates
 - Real-time depending on N

Visual tracking

Standard tracking algorithms

- **Motion segmentation based tracking**
- **Template tracking**
- **Mean-shift tracking**
- **Kalman filter**
- **Particle filtering**

Particle filtering



- Start with weighted samples from previous time step
- Sample and shift according to dynamics model
- Spread due to randomness; this is predicted density $p(x_t|y_{t-1})$
- Weight the samples according to observation density
- Arrive at corrected density estimate $p(x_t|y_t)$